

20104169 - SUMESH R

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
from google.colab import drive
drive.mount('/content/drive')
df=pd.read_csv("/content/drive/MyDrive/mydatasets/csvs_per_year/madrid_2007.csv")
df
```

Mounted at /content/drive

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2
0	2007-12-01 01:00:00	NaN	2.86	NaN	NaN	NaN	282.200012	1054.000000	NaN	4.030000	156.199997	97.43	NaN	64.519999
1	2007-12-01 01:00:00	NaN	1.82	NaN	NaN	NaN	86.419998	354.600006	NaN	3.260000	80.809998	NaN	NaN	35.419999
2	2007-12-01 01:00:00	NaN	1.47	NaN	NaN	NaN	94.639999	319.000000	NaN	5.310000	53.099998	NaN	NaN	19.080000
3	2007-12-01 01:00:00	NaN	1.64	NaN	NaN	NaN	127.900002	476.700012	NaN	4.500000	105.300003	NaN	NaN	17.670000
4	2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.500000	15.90	3.56	40.230000
...
225115	2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	6.760000	5.14	1.00	7.420000
225116	2007-03-01 00:00:00	NaN	0.16	NaN	NaN	NaN	46.820000	51.480000	NaN	22.150000	5.700000	NaN	NaN	7.130000
225117	2007-03-01 00:00:00	0.24	NaN	0.20	NaN	0.09	51.259998	66.809998	NaN	18.540001	13.010000	6.95	NaN	8.740000
225118	2007-03-01 00:00:00	0.11	NaN	1.00	NaN	0.05	24.240000	36.930000	NaN	NaN	6.610000	NaN	NaN	9.890000
225119	2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	10.260000	7.08	1.23	9.890000

225120 rows x 17 columns



Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        25443 non-null  object
 1   BEN         25443 non-null  float64
 2   CO          25443 non-null  float64
 3   EBE         25443 non-null  float64
 4   MXY         25443 non-null  float64
 5   NMHC        25443 non-null  float64
 6   NO_2        25443 non-null  float64
 7   NOx         25443 non-null  float64
 8   OXY         25443 non-null  float64
 9   O_3         25443 non-null  float64
10  PM10        25443 non-null  float64
11  PM25        25443 non-null  float64
12  PXY         25443 non-null  float64
13  SO_2        25443 non-null  float64
14  TCH         25443 non-null  float64
15  TOL         25443 non-null  float64
16  station     25443 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [6]:

```
data=df[['CO' , 'station']]
data
```

Out[6]:

	CO	station
4	1.86	28079006
21	0.31	28079024
25	1.42	28079099
30	1.89	28079006
47	0.30	28079024
...
225073	0.47	28079006
225094	0.45	28079099
225098	0.41	28079006
225115	0.45	28079024
225119	0.40	28079099


Line chart

In [7]:

```
data.plot.line(subplots=True)
```

Out[7]:

array([<Axes: >, <Axes: >], dtype=object)



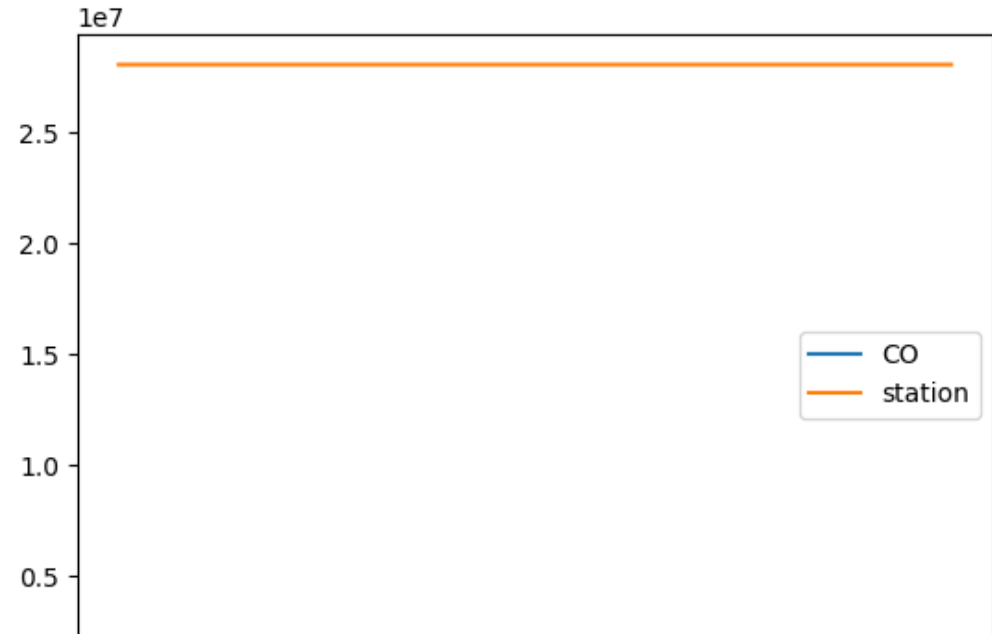
Line chart


In [8]:

```
data.plot.line()
```

Out[8]:

<Axes: >





Bar chart

In [9]:


```
b=data[0:50]
```

In [10]:

```
b.plot.bar()
```

Out[10]:

<Axes: >




Histogram


In [11]:

```
data.plot.hist()
```

Out[11]:

<Axes: ylabel='Frequency'>






Area chart

In [12]:

```
data.plot.area()
```

Out[12]:

<Axes: >



Box chart


In [13]:

```
data.plot.box()
```

Out[13]:

<Axes: >






Pie chart


In [14]:

```
b.plot.pie(y='station' )
```

Out[14]:

<Axes: ylabel='station'>






Scatter chart

In [15]:

```
data.plot.scatter(x='CO', y='station')
```

Out[15]:

<Axes: xlabel='CO', ylabel='station'>



In [16]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   date        25443 non-null  object
1   PM2.5       25443 non-null  float64
```

```
1 BEN 25443 non-null float64
2 CO 25443 non-null float64
3 EBE 25443 non-null float64
4 MXY 25443 non-null float64
5 NMHC 25443 non-null float64
6 NO_2 25443 non-null float64
7 NOx 25443 non-null float64
8 OXY 25443 non-null float64
9 O_3 25443 non-null float64
10 PM10 25443 non-null float64
11 PM25 25443 non-null float64
12 PXY 25443 non-null float64
13 SO_2 25443 non-null float64
14 TCH 25443 non-null float64
15 TOL 25443 non-null float64
16 station 25443 non-null int64
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [17]:

```
df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	
count	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	254
mean	1.146744	0.505120	1.394071	2.392008	0.249967	58.532683	112.741861	1.270278	
std	1.278733	0.423231	1.268265	2.784302	0.142627	37.755029	115.527006	1.143188	
min	0.130000	0.000000	0.120000	0.150000	0.000000	1.690000	1.780000	0.110000	
25%	0.450000	0.260000	0.780000	0.960000	0.160000	31.285001	39.910000	0.740000	
50%	0.770000	0.400000	1.000000	1.500000	0.220000	54.080002	82.809998	1.000000	
75%	1.390000	0.640000	1.580000	2.855000	0.300000	79.230003	149.300003	1.360000	
max	30.139999	9.660000	31.680000	65.480003	2.570000	430.299988	1893.000000	32.540001	1

In [18]:

```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
        'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

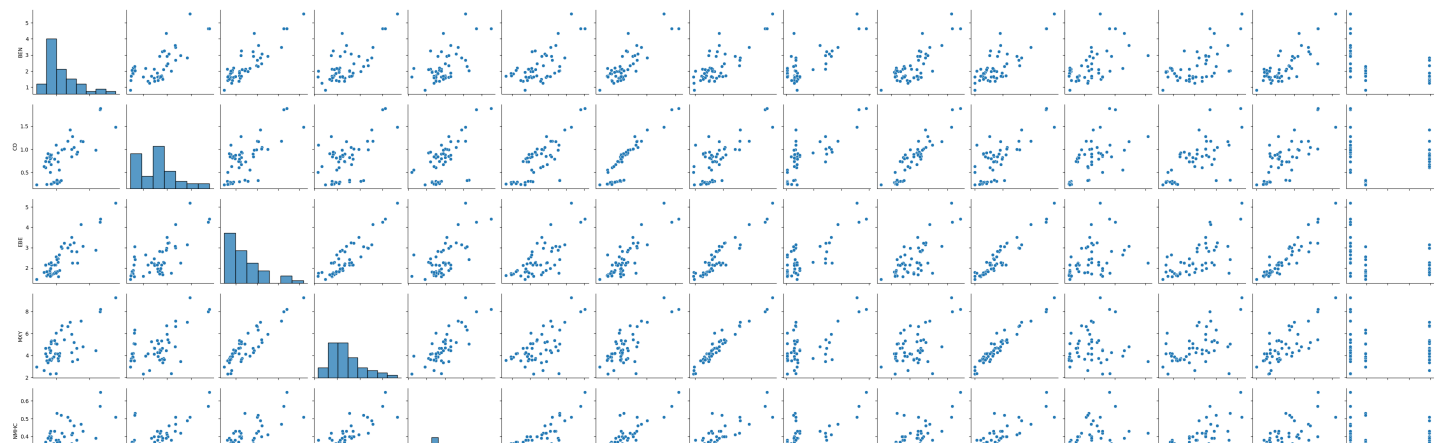
EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

Out[19]:

<seaborn.axisgrid.PairGrid at 0x7c28d80a6080>





In [20]:

```
sns.distplot(df1['station'])
```

<ipython-input-20-4bc330f7257f>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df1['station'])
```

Out[20]:

<Axes: xlabel='station', ylabel='Density'>






In [21]:

```
sns.heatmap(df1.corr())
```

Out[21]:

<Axes: >



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

In [23]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

In [24]:

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[24]:

```
LinearRegression
```

In [25]:

```
lr.intercept_
```

Out[25]:

28079009.537545096

In [26]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[26]:


Co-efficient	
BEN	-33.878090
CO	16.498366
EBE	1.482030
MXY	-1.353057
NMHC	-40.412680
NO_2	0.121295
NOx	-0.037018
OXY	6.095063
O_3	-0.036906
PM10	0.143900
PXY	7.256940
SO_2	0.246843
TCH	25.289281
TOL	3.010203


In [27]:

```
prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[27]:

<matplotlib.collections.PathCollection at 0x7c28c503e3e0>





ACCURACY

In [28]:

```
lr.score(x_test,y_test)
```

Out[28]:

0.15525827694238103

In [29]:

```
lr.score(x_train,y_train)
```

Out[29]:

0.16056031664547565

Ridge and Lasso

In [30]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [31]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[31]:

▼ Ridge
Ridge(alpha=10)

Accuracy(Ridge)

In [32]:

```
rr.score(x_test,y_test)
```

Out[32]:

0.1552057727588495

In [33]:

```
rr.score(x_train,y_train)
```

Out[33]:

0.16050892632951208

In [34]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[34]:

▼ Lasso
Lasso(alpha=10)

```
Lasso(alpha=10)
```

In [35]:

```
la.score(x_train,y_train)
```

Out[35]:

```
0.01325224953187909
```

Accuracy(Lasso)

In [36]:

```
la.score(x_test,y_test)
```

Out[36]:

```
0.014198799959927011
```

In [37]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]:

```
▼ ElasticNet
ElasticNet()
```

In [38]:

```
en.coef_
```

Out[38]:

```
array([-7.99626026,  0.          ,  0.          ,  0.21923228, -0.          ,
        0.06336065, -0.0554393 ,  0.8249646 , -0.05906873,  0.16760411,
        0.72498793,  0.01837265,  0.          ,  0.84622624])
```

In [39]:

```
en.intercept_
```

Out[39]:

```
28079045.564106826
```

In [40]:

```
prediction=en.predict(x_test)
```

In [41]:

```
en.score(x_test,y_test)
```

Out[41]:

```
0.06927782673837335
```

Evaluation Metrics

In [42]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
```

```
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
36.70207687149274
1535.6845903561043
39.18781175768946
```

Logistic Regression

In [43]:

```
from sklearn.linear_model import LogisticRegression
```

In [44]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

In [45]:

```
feature_matrix.shape
```

Out[45]:

```
(25443, 14)
```

In [46]:

```
target_vector.shape
```

Out[46]:

```
(25443,)
```

In [47]:

```
from sklearn.preprocessing import StandardScaler
```

In [48]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [49]:

```
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

Out[49]:

```
▼      LogisticRegression
LogisticRegression(max_iter=10000)
```

In [50]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

In [51]:

```
prediction=logr.predict(observation)
print(prediction)
```

```
[28079099]
```

In [52]:

```
logr.classes_
```

Out[52]:

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14])
```

```
array([28079006, 28079024, 28079099])
```

In [53]:

```
logr.score(fs,target_vector)
```

Out[53]:

```
0.8146838030106512
```

In [54]:

```
logr.predict_proba(observation)[0][0]
```

Out[54]:

```
1.0827539764163807e-19
```

In [55]:

```
logr.predict_proba(observation)
```

Out[55]:

```
array([[1.08275398e-19, 1.80383815e-19, 1.00000000e+00]])
```

Random Forest

In [56]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [57]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[57]:

```
▼ RandomForestClassifier  
RandomForestClassifier()
```

In [58]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
}
```

In [59]:

```
from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[59]:

```
┌────────────────── GridSearchCV ───────────────────┐  
│ ▶ estimator: RandomForestClassifier                  │  
│ ┌────────────────── RandomForestClassifier ─────────┐ │  
│ │                                                    │ │  
│ └──────────────────┘ │  
└──────────────────┘
```

In [60]:

```
grid_search.best_score_
```

Out[60]:

```
0.8214486243683324
```

In [61]:

```
rfc_best=grid_search.best_estimator_
```

In [62]:

```
from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

Out[62]:

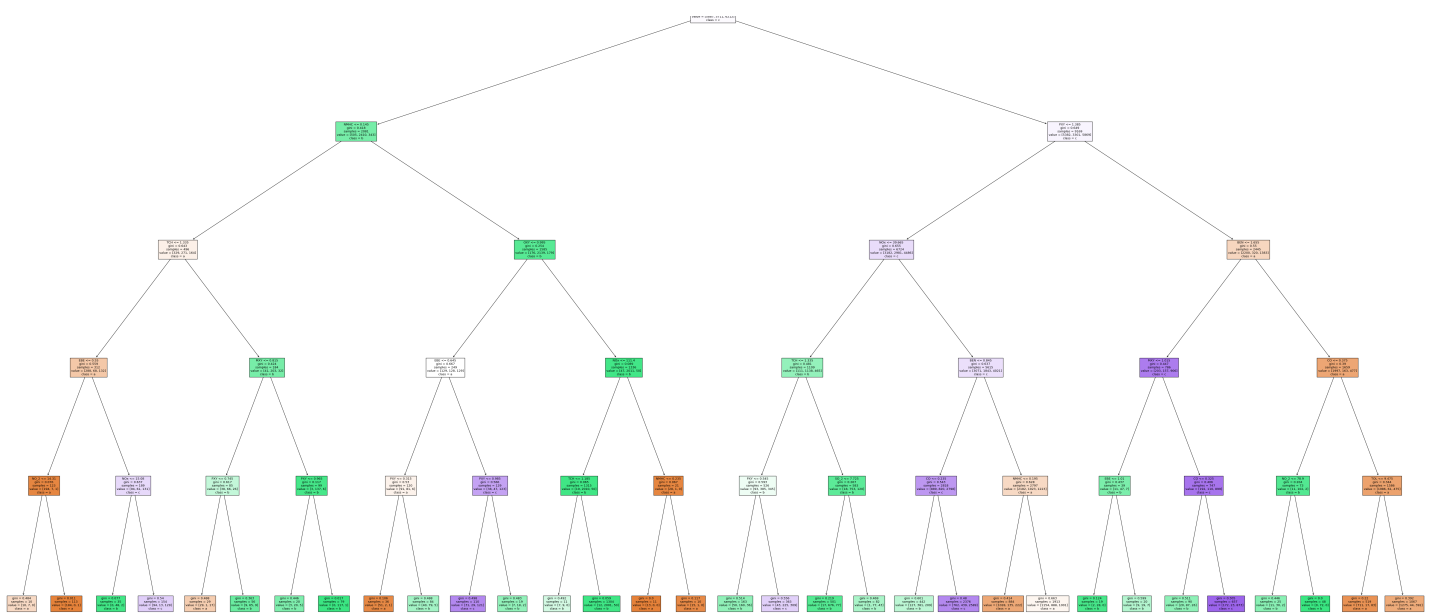
```
[Text(0.5, 0.9166666666666666, 'TOL <= 1.325\ngini = 0.666\nsamples = 11250\nvalue = [588
7, 5711, 6212]\nlass = c'),
 Text(0.25, 0.75, 'NMHC <= 0.145\ngini = 0.418\nsamples = 2081\nvalue = [505, 2410, 343]\nlass = b'),
 Text(0.125, 0.5833333333333333, 'TCH <= 1.335\ngini = 0.643\nsamples = 496\nvalue = [329
, 271, 164]\nlass = a'),
 Text(0.0625, 0.4166666666666667, 'EBE <= 0.55\ngini = 0.559\nsamples = 312\nvalue = [288
, 68, 132]\nlass = a'),
 Text(0.03125, 0.25, 'NO_2 <= 14.31\ngini = 0.076\nsamples = 123\nvalue = [194, 7, 1]\nlass = a'),
 Text(0.015625, 0.0833333333333333, 'gini = 0.484\nsamples = 10\nvalue = [10, 7, 0]\nlass = a'),
 Text(0.046875, 0.0833333333333333, 'gini = 0.011\nsamples = 113\nvalue = [184, 0, 1]\nlass = a'),
 Text(0.09375, 0.25, 'NOx <= 15.08\ngini = 0.637\nsamples = 189\nvalue = [94, 61, 131]\nlass = c'),
 Text(0.078125, 0.0833333333333333, 'gini = 0.077\nsamples = 35\nvalue = [0, 48, 2]\nlass = b'),
 Text(0.109375, 0.0833333333333333, 'gini = 0.54\nsamples = 154\nvalue = [94, 13, 129]\nlass = c'),
 Text(0.1875, 0.4166666666666667, 'MXY <= 0.815\ngini = 0.424\nsamples = 184\nvalue = [41
, 203, 32]\nlass = b'),
 Text(0.15625, 0.25, 'PXY <= 0.745\ngini = 0.617\nsamples = 85\nvalue = [38, 66, 26]\nlass = b'),
 Text(0.140625, 0.0833333333333333, 'gini = 0.488\nsamples = 29\nvalue = [29, 1, 17]\nlass = a'),
 Text(0.171875, 0.0833333333333333, 'gini = 0.363\nsamples = 56\nvalue = [9, 65, 9]\nlass = b'),
 Text(0.21875, 0.25, 'PXY <= 0.965\ngini = 0.117\nsamples = 99\nvalue = [3, 137, 6]\nlass = b'),
 Text(0.203125, 0.0833333333333333, 'gini = 0.446\nsamples = 20\nvalue = [3, 20, 5]\nlass = b'),
 Text(0.234375, 0.0833333333333333, 'gini = 0.017\nsamples = 79\nvalue = [0, 117, 1]\nlass = b'),
 Text(0.375, 0.5833333333333333, 'OXY <= 0.995\ngini = 0.254\nsamples = 1585\nvalue = [17
6, 2139, 179]\nlass = b'),
 Text(0.3125, 0.4166666666666667, 'EBE <= 0.645\ngini = 0.667\nsamples = 249\nvalue = [12
9, 128, 129]\nlass = a'),
 Text(0.28125, 0.25, 'PXY <= 0.315\ngini = 0.53\nsamples = 120\nvalue = [91, 81, 6]\nlass = a'),
 Text(0.265625, 0.0833333333333333, 'gini = 0.106\nsamples = 36\nvalue = [51, 2, 1]\nlass = a'),
 Text(0.296875, 0.0833333333333333, 'gini = 0.488\nsamples = 84\nvalue = [40, 79, 5]\nlass = b'),
 Text(0.34375, 0.25, 'PXY <= 0.985\ngini = 0.566\nsamples = 129\nvalue = [38, 47, 123]\nlass = c'),
 Text(0.328125, 0.0833333333333333, 'gini = 0.498\nsamples = 110\nvalue = [31, 29, 121]\nlass = c'),
 Text(0.359375, 0.0833333333333333, 'gini = 0.483\nsamples = 19\nvalue = [7, 18, 2]\nlass = b'),
 Text(0.4375, 0.4166666666666667, 'NOx <= 111.4\ngini = 0.089\nsamples = 1336\nvalue = [4
7, 2011, 50]\nlass = b'),
 Text(0.40625, 0.25, 'TCH <= 1.185\ngini = 0.065\nsamples = 1315\nvalue = [19, 2010, 50]\nlass = b'),
 Text(0.390625, 0.0833333333333333, 'gini = 0.492\nsamples = 11\nvalue = [7, 9, 0]\nlass = b'),
 Text(0.421875, 0.0833333333333333, 'gini = 0.059\nsamples = 1304\nvalue = [12, 2001, 50
```



```

]\nclasse = b'),
  Text(0.46875, 0.25, 'NMHC <= 0.235\ngini = 0.067\nsamples = 21\nvalue = [28, 1, 0]\nclasse = a'),
  Text(0.453125, 0.08333333333333333, 'gini = 0.0\nsamples = 11\nvalue = [13, 0, 0]\nclasse = a'),
  Text(0.484375, 0.08333333333333333, 'gini = 0.117\nsamples = 10\nvalue = [15, 1, 0]\nclasse = a'),
  Text(0.75, 0.75, 'PXY <= 1.385\ngini = 0.649\nsamples = 9169\nvalue = [5382, 3301, 5869]\nclasse = c'),
  Text(0.625, 0.5833333333333334, 'NOx <= 39.665\ngini = 0.655\nsamples = 6724\nvalue = [3182, 2981, 4486]\nclasse = c'),
  Text(0.5625, 0.4166666666666667, 'TCH <= 1.335\ngini = 0.481\nsamples = 1109\nvalue = [111, 1138, 465]\nclasse = b'),
  Text(0.53125, 0.25, 'PXY <= 0.545\ngini = 0.593\nsamples = 526\nvalue = [93, 385, 345]\nclasse = b'),
  Text(0.515625, 0.08333333333333333, 'gini = 0.514\nsamples = 163\nvalue = [50, 160, 36]\nclasse = b'),
  Text(0.546875, 0.08333333333333333, 'gini = 0.556\nsamples = 363\nvalue = [43, 225, 309]\nclasse = c'),
  Text(0.59375, 0.25, 'SO_2 <= 7.725\ngini = 0.267\nsamples = 583\nvalue = [18, 753, 120]\nclasse = b'),
  Text(0.578125, 0.08333333333333333, 'gini = 0.219\nsamples = 501\nvalue = [17, 676, 77]\nclasse = b'),
  Text(0.609375, 0.08333333333333333, 'gini = 0.469\nsamples = 82\nvalue = [1, 77, 43]\nclasse = b'),
  Text(0.6875, 0.4166666666666667, 'BEN <= 0.845\ngini = 0.637\nsamples = 5615\nvalue = [3071, 1843, 4021]\nclasse = c'),
  Text(0.65625, 0.25, 'CO <= 0.235\ngini = 0.543\nsamples = 2818\nvalue = [889, 820, 2798]\nclasse = c'),
  Text(0.640625, 0.08333333333333333, 'gini = 0.601\nsamples = 442\nvalue = [127, 381, 209]\nclasse = b'),
  Text(0.671875, 0.08333333333333333, 'gini = 0.48\nsamples = 2376\nvalue = [762, 439, 2589]\nclasse = c'),
  Text(0.71875, 0.25, 'NMHC <= 0.195\ngini = 0.628\nsamples = 2797\nvalue = [2182, 1023, 1223]\nclasse = a'),
  Text(0.703125, 0.08333333333333333, 'gini = 0.414\nsamples = 884\nvalue = [1028, 135, 222]\nclasse = a'),
  Text(0.734375, 0.08333333333333333, 'gini = 0.663\nsamples = 1913\nvalue = [1154, 888, 1001]\nclasse = a'),
  Text(0.875, 0.5833333333333334, 'BEN <= 1.655\ngini = 0.55\nsamples = 2445\nvalue = [2200, 320, 1383]\nclasse = a'),
  Text(0.8125, 0.4166666666666667, 'MXY <= 1.015\ngini = 0.447\nsamples = 786\nvalue = [203, 157, 906]\nclasse = c'),
  Text(0.78125, 0.25, 'EBE <= 1.01\ngini = 0.437\nsamples = 39\nvalue = [11, 47, 7]\nclasse = b'),
  Text(0.765625, 0.08333333333333333, 'gini = 0.124\nsamples = 19\nvalue = [2, 28, 0]\nclasse = b'),
  Text(0.796875, 0.08333333333333333, 'gini = 0.599\nsamples = 20\nvalue = [9, 19, 7]\nclasse = b'),
  Text(0.84375, 0.25, 'CO <= 0.325\ngini = 0.406\nsamples = 747\nvalue = [192, 110, 899]\nclasse = c'),
  Text(0.828125, 0.08333333333333333, 'gini = 0.511\nsamples = 90\nvalue = [20, 87, 26]\nclasse = b'),
  Text(0.859375, 0.08333333333333333, 'gini = 0.305\nsamples = 657\nvalue = [172, 23, 873]\nclasse = c'),
  Text(0.9375, 0.4166666666666667, 'CO <= 0.375\ngini = 0.39\nsamples = 1659\nvalue = [1997, 163, 477]\nclasse = a'),
  Text(0.90625, 0.25, 'NO_2 <= 78.9\ngini = 0.204\nsamples = 73\nvalue = [11, 102, 2]\nclasse = b'),
  Text(0.890625, 0.08333333333333333, 'gini = 0.446\nsamples = 25\nvalue = [11, 30, 2]\nclasse = b'),
  Text(0.921875, 0.08333333333333333, 'gini = 0.0\nsamples = 48\nvalue = [0, 72, 0]\nclasse = b'),
  Text(0.96875, 0.25, 'TOL <= 9.475\ngini = 0.344\nsamples = 1586\nvalue = [1986, 61, 475]\nclasse = a'),
  Text(0.953125, 0.08333333333333333, 'gini = 0.22\nsamples = 519\nvalue = [711, 17, 83]\nclasse = a'),
  Text(0.984375, 0.08333333333333333, 'gini = 0.392\nsamples = 1067\nvalue = [1275, 44, 392]\nclasse = a')]

```



Conclusion

Accuracy

In [63]:

```
print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

Linear Regression: 0.15525827694238103
Ridge Regression: 0.1552057727588495
Lasso Regression 0.014198799959927011
ElasticNet Regression: 0.06927782673837335
Logistic Regression: 0.8146838030106512
Random Forest: 0.8214486243683324

Random Forest is suitable for this dataset