

# 20104169 - SUMESH R

## Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
from google.colab import drive
drive.mount('/content/drive')
df=pd.read_csv("/content/drive/MyDrive/mydatasets/csvs_per_year/madrid_2013.csv")
df
```

Mounted at /content/drive

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	NaN	0.6	NaN	NaN	135.0	74.0	NaN	NaN	NaN	7.0	NaN	NaN	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	NaN	71.0	83.0	2.0	23.0	16.0	12.0	NaN	8.3	28079008
2	2013-11-01 01:00:00	3.9	NaN	2.8	NaN	49.0	70.0	NaN	NaN	NaN	NaN	NaN	9.0	28079011
3	2013-11-01 01:00:00	NaN	0.5	NaN	NaN	82.0	87.0	3.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2013-11-01 01:00:00	NaN	NaN	NaN	NaN	242.0	111.0	2.0	NaN	NaN	12.0	NaN	NaN	28079017
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209875	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	8.0	39.0	52.0	NaN	NaN	NaN	NaN	NaN	28079056
209876	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	1.0	11.0	NaN	6.0	NaN	2.0	NaN	NaN	28079057
209877	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	4.0	75.0	NaN	NaN	NaN	NaN	NaN	28079058
209878	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	11.0	52.0	NaN	NaN	NaN	NaN	NaN	28079059
209879	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	1.0	10.0	75.0	3.0	NaN	NaN	NaN	NaN	28079060

209880 rows × 14 columns

## Data Cleaning and Data Preprocessing

In [3]:

```
df=df.fillna(1)
df
```

Out[3]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2013-11-01 01:00:00	1.0	0.6	1.0	1.0	135.0	74.0	1.0	1.0	1.0	7.0	1.0	1.0	28079004
1	2013-11-01 01:00:00	1.5	0.5	1.3	1.0	71.0	83.0	2.0	23.0	16.0	12.0	1.0	8.3	28079008
2	2013-11-01 01:00:00	3.9	1.0	2.8	1.0	49.0	70.0	1.0	1.0	1.0	1.0	1.0	9.0	28079011
3	2013-11-01 01:00:00	1.0	0.5	1.0	1.0	82.0	87.0	3.0	1.0	1.0	1.0	1.0	1.0	28079016
4	2013-11-01 01:00:00	1.0	1.0	1.0	1.0	242.0	111.0	2.0	1.0	1.0	12.0	1.0	1.0	28079017

...	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
209875	2013-03-01 00:00:00	1.0	0.4	1.0	1.0	8.0	39.0	52.0	1.0	1.0	1.0	1.0	1.0	28079056
209876	2013-03-01 00:00:00	1.0	0.4	1.0	1.0	1.0	11.0	1.0	6.0	1.0	2.0	1.0	1.0	28079057
209877	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	2.0	4.0	75.0	1.0	1.0	1.0	1.0	1.0	28079058
209878	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	2.0	11.0	52.0	1.0	1.0	1.0	1.0	1.0	28079059
209879	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	1.0	10.0	75.0	3.0	1.0	1.0	1.0	1.0	28079060

209880 rows x 14 columns

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
      'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209880 non-null  object
1   BEN         209880 non-null  float64
2   CO          209880 non-null  float64
3   EBE         209880 non-null  float64
4   NMHC        209880 non-null  float64
5   NO          209880 non-null  float64
6   NO_2        209880 non-null  float64
7   O_3         209880 non-null  float64
8   PM10        209880 non-null  float64
9   PM25        209880 non-null  float64
10  SO_2        209880 non-null  float64
11  TCH         209880 non-null  float64
12  TOL         209880 non-null  float64
13  station     209880 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

In [6]:

```
data=df[['CO' , 'station']]
data
```

Out[6]:

	CO	station
0	0.6	28079004
1	0.5	28079008
2	1.0	28079011
3	0.5	28079016
4	1.0	28079017
...	...	...
209875	0.4	28079056
209876	0.4	28079057
209877	1.0	28079058
209878	1.0	28079059

209880 rows x 2 columns

209880 rows x 2 columns


## Line chart

In [7]:

```
data.plot.line(subplots=True)
```

Out[7]:

array([<Axes: >, <Axes: >], dtype=object)



## Line chart


In [8]:

```
data.plot.line()
```

Out[8]:

<Axes: >





# Bar chart

In [9]:


```
b=data[0:50]
```

In [10]:

```
b.plot.bar()
```

Out[10]:

<Axes: >




# Histogram

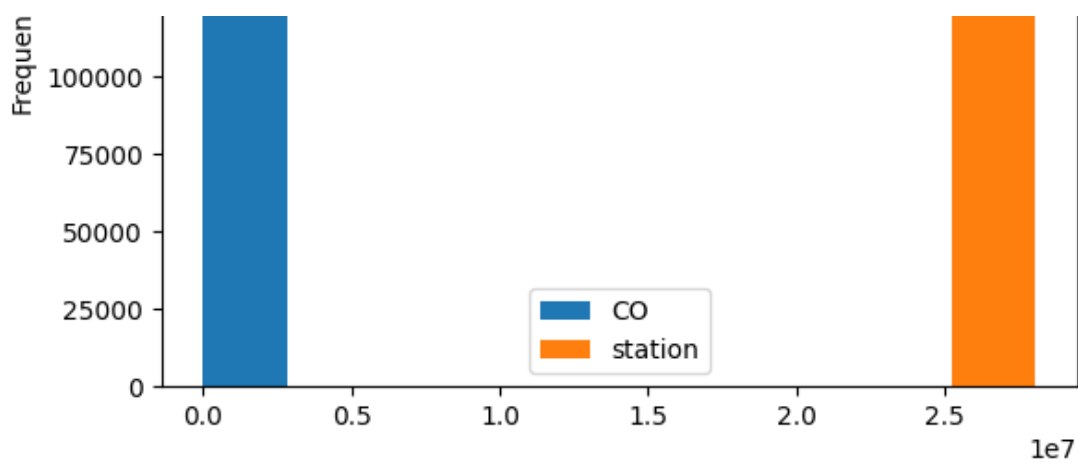
In [11]:

```
data.plot.hist()
```

Out[11]:

<Axes: ylabel='Frequency'>





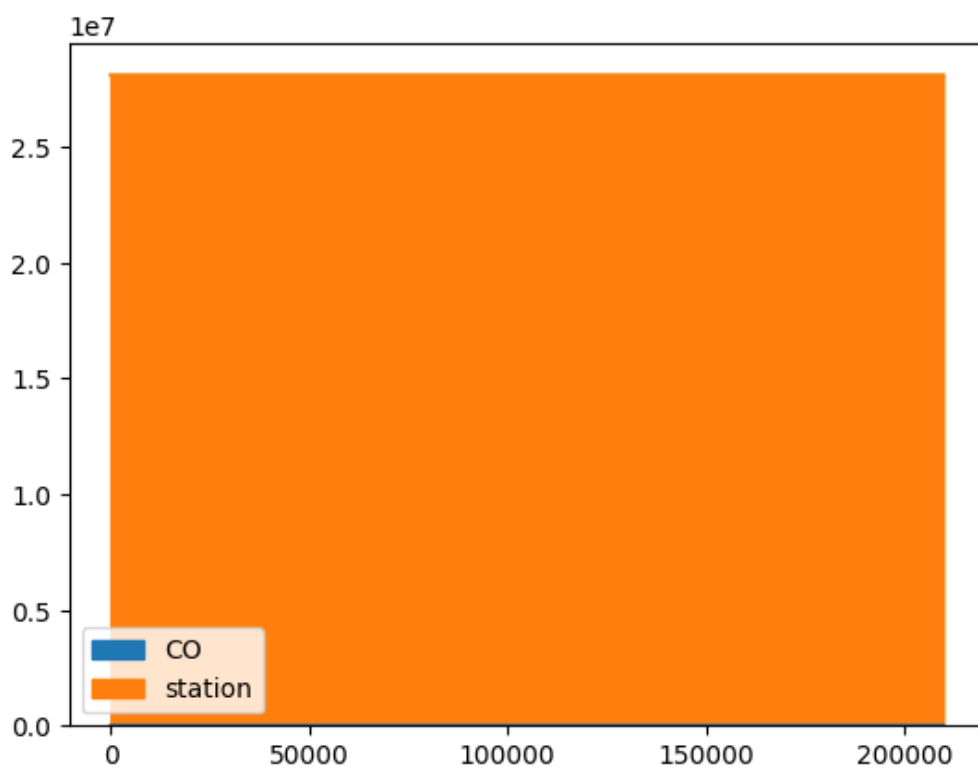
## Area chart

In [12]:

```
data.plot.area()
```

Out[12]:

<Axes: >



## Box chart


In [13]:

```
data.plot.box()
```

Out[13]:

<Axes: >





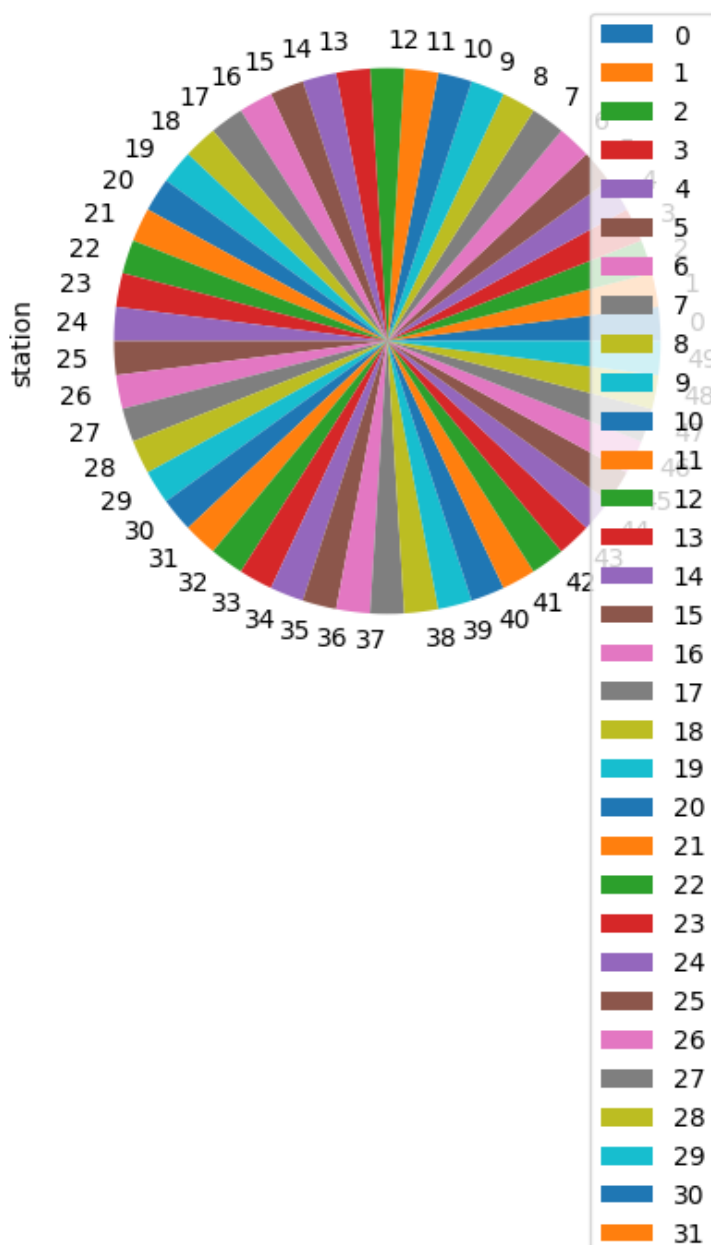
## Pie chart


In [14]:

```
b.plot.pie(y='station' )
```

Out[14]:

<Axes: ylabel='station'>






## Scatter chart

In [15]:

```
data.plot.scatter(x='CO' ,y='station')
```

Out[15]:

<Axes: xlabel='CO', ylabel='station'>



In [16]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column  Non-Null Count  Dtype
---

```

```
0    date      209880 non-null object
1    BEN       209880 non-null float64
2    CO        209880 non-null float64
3    EBE       209880 non-null float64
4    NMHC      209880 non-null float64
5    NO        209880 non-null float64
6    NO_2      209880 non-null float64
7    O_3       209880 non-null float64
8    PM10      209880 non-null float64
9    PM25      209880 non-null float64
10   SO_2      209880 non-null float64
11   TCH       209880 non-null float64
12   TOL       209880 non-null float64
13   station   209880 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

In [17]:

```
df.describe()
```

Out[17]:

	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PI
count	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000
mean	0.931014	0.721695	0.954744	0.900223	20.101401	34.586402	29.461235	9.636
std	0.430684	0.361528	0.301074	0.267139	44.319112	27.866588	35.362880	13.492
min	0.100000	0.100000	0.100000	0.040000	1.000000	1.000000	1.000000	1.000
25%	1.000000	0.300000	1.000000	1.000000	2.000000	14.000000	1.000000	1.000
50%	1.000000	1.000000	1.000000	1.000000	5.000000	27.000000	8.000000	1.000
75%	1.000000	1.000000	1.000000	1.000000	17.000000	48.000000	54.000000	14.000
max	12.100000	10.400000	11.800000	1.000000	1081.000000	388.000000	226.000000	232.000


## EDA AND VISUALIZATION

In [18]:


```
sns.pairplot(df[0:50])
```

Out[18]:

<seaborn.axisgrid.PairGrid at 0x7d1a405fcbb0>







In [19]:

```
sns.distplot(df['station'])
```

<ipython-input-19-6e2460d4583e>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['station'])
```

Out[19]:

<Axes: xlabel='station', ylabel='Density'>



In [20]:


```
sns.heatmap(df.corr())
```

<ipython-input-20-aa4f4450a243>:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
sns.heatmap(df.corr())
```

Out[20]:

<Axes: >



## TO TRAIN THE MODEL AND MODEL BUILDING

In [21]:

```
x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

In [22]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

In [23]:

```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[23]:

```
▼ LinearRegression  
LinearRegression()
```

In [24]:

```
lr.intercept_
```

Out[24]:

28078977.103350364

In [25]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[25]:


Co-efficient	
BEN	1.675295
CO	18.812592
EBE	10.008459
NMHC	18.092884
NO	-0.009336
NO_2	-0.039330
O_3	0.008942
PM10	0.279728
PM25	-0.371620
SO_2	-0.945909
TCH	25.161095
TOL	-3.445599

In [26]:

```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[26]:

<matplotlib.collections.PathCollection at 0x7d1a7f509660>



# ACCURACY

In [27]:

```
lr.score(x_test,y_test)
```

Out[27]:

```
0.30543236748624947
```

In [28]:

```
lr.score(x_train,y_train)
```

Out[28]:

```
0.30799712227942355
```

## Ridge and Lasso

In [29]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [30]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[30]:

```
▼      Ridge  
Ridge(alpha=10)
```

## Accuracy(Ridge)

In [31]:

```
rr.score(x_test,y_test)
```

Out[31]:

```
0.3054128646053008
```

In [32]:

```
rr.score(x_train,y_train)
```

Out[32]:

```
0.3079943077378543
```

In [33]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[33]:

```
▼      Lasso  
Lasso(alpha=10)
```

In [34]:

```
la.score(x_train,y_train)
```

Out[34]:

```
0.0451475937957041
```

## Accuracy(Lasso)

In [35]:

```
la.score(x_test,y_test)
```

Out[35]:

```
0.04531888865719813
```

In [36]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[36]:

```
▼ ElasticNet
ElasticNet()
```

In [37]:

```
en.coef_
```

Out[37]:

```
array([ 0.33177767,  2.64130733,  0.45537669,  0.          ,  0.03569569,
        -0.05608388, -0.01983977,  0.23701462, -0.35214909, -1.35812245,
        -0.          , -1.5728494 ])
```

In [38]:

```
en.intercept_
```

Out[38]:

```
28079041.198403195
```

In [39]:

```
prediction=en.predict(x_test)
```

In [40]:

```
en.score(x_test,y_test)
```

Out[40]:

```
0.1609692866422956
```

## Evaluation Metrics

In [41]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
13.81701650598825
259.6923434931426
16.11497264946927
```

# Logistic Regression

In [42]:

```
from sklearn.linear_model import LogisticRegression
```

In [43]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
                  'SO_2', 'TCH', 'TOL']][0:50]  
target_vector=df[ 'station'][0:50]
```

In [44]:

```
feature_matrix.shape
```

Out[44]:

```
(50, 12)
```

In [45]:

```
target_vector.shape
```

Out[45]:

```
(50,)
```

In [46]:

```
from sklearn.preprocessing import StandardScaler
```

In [47]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [48]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

Out[48]:

```
▼ LogisticRegression  
LogisticRegression(max_iter=10000)
```

In [49]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
```

In [50]:

```
prediction=logr.predict(observation)  
print(prediction)
```

```
[28079038]
```

In [51]:

```
logr.classes_
```

Out[51]:

```
array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,  
       28079024, 28079027, 28079035, 28079036, 28079038, 28079039,  
       28079040, 28079047, 28079048, 28079049, 28079050, 28079054,  
       28079055, 28079056, 28079057, 28079058, 28079059, 28079060])
```

```
logr.score(fs, target_vector)
```

0.9

```
logr.predict_proba(observation)[0][0]
```

3.2424612943857784e-12

```
logr.predict proba(observation)
```

```
array([[3.24246129e-12, 1.98879790e-01, 2.68546600e-11, 9.18595327e-17,
        2.85595544e-05, 1.13436010e-07, 1.11387803e-06, 1.83744227e-14,
        4.41788720e-09, 2.60721644e-13, 8.01089957e-01, 6.75335556e-16,
        4.31028393e-12, 3.50004736e-07, 1.45256021e-14, 1.06842862e-19,
        1.45341132e-13, 1.18084324e-14, 1.11311169e-07, 4.24582423e-10,
        4.14877181e-17, 1.07641028e-18, 8.68787036e-11, 5.39240034e-11]])
```

## In [55]:

```
from sklearn.ensemble import RandomForestClassifier
```

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
▼ RandomForestClassifier
RandomForestClassifier()
```

```
parameters={'max_depth':[1,2,3,4,5],
            'min_samples_leaf':[5,10,15,20,25],
            'n_estimators':[10,20,30,40,50]
}
```

```
from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

```

GridSearchCV
├── estimator: RandomForestClassifier
│   └── RandomForestClassifier
└── ...

```

```
grid_search.best_score
```

Out[59]:

0.7496052165863487

In [60]:

```
rfc_best=grid_search.best_estimator_
```

In [61]:

```
from sklearn.tree import plot_tree

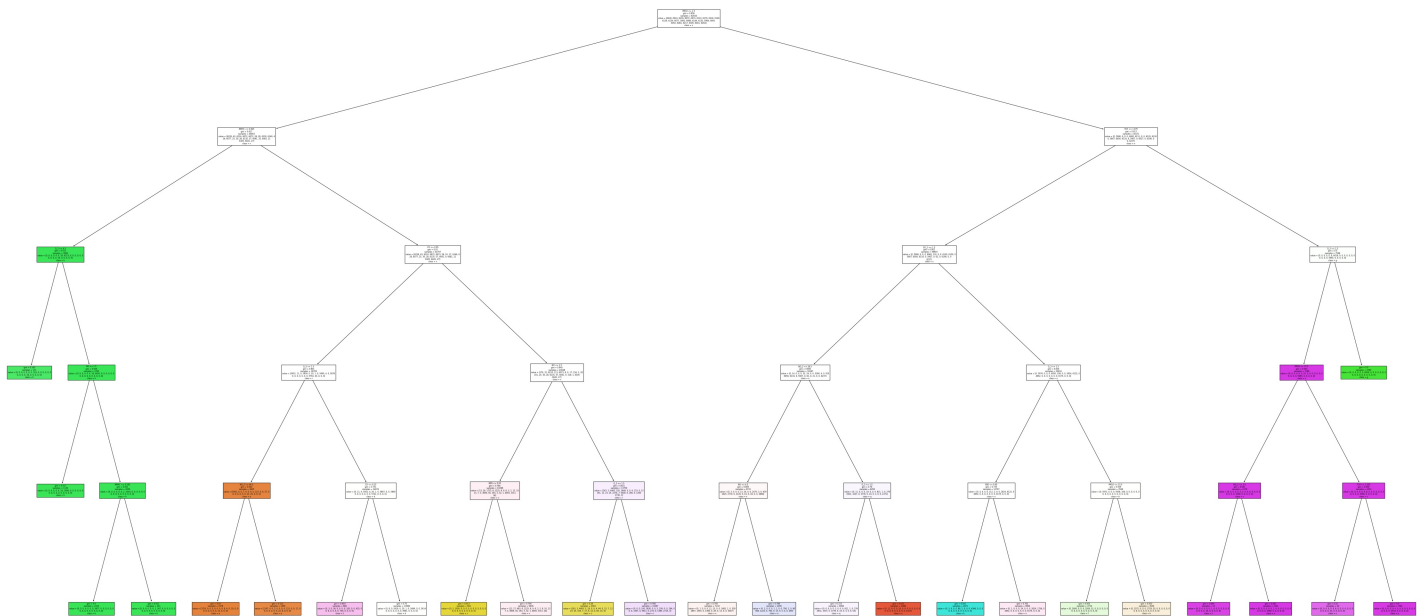
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d','e',
'f','g','h','i','j','k','l','m','n','o','p','q','r','s','t','u','v','w','x','y','z'],fi
lled=True)
```

Out[61]:

```
[Text(0.483695652173913, 0.9166666666666666, 'PM10 <= 1.5\ngini = 0.958\nsamples = 92834\nvalue = [6228, 6033, 6216, 6072, 6073, 6110, 6179, 6230, 6146\n6126, 6154, 6077, 5892, 6086, 6144, 6135, 5964, 6091\n6052, 6081, 6217, 6305, 6041, 6264]\nlass = v'),
Text(0.17391304347826086, 0.75, 'NMHC <= 0.845\ngini = 0.917\nsamples = 46663\nvalue = [6228, 43, 6216, 6072, 6073, 28, 28, 6230, 6146, 6\n28, 6077, 25, 30, 28, 6135, 57, 6091, 25, 6081, 21\n6305, 6041, 27]\nlass = v'),
Text(0.043478260869565216, 0.5833333333333334, 'O_3 <= 4.5\ngini = 0.011\nsamples = 3906\nvalue = [0, 0, 0, 0, 0, 0, 18, 6213, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 16, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.021739130434782608, 0.4166666666666667, 'gini = 0.163\nsamples = 111\nvalue = [0, 0, 0, 0, 0, 0, 0, 163, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 16, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.06521739130434782, 0.4166666666666667, 'NO <= 1.5\ngini = 0.006\nsamples = 3795\nvalue = [0, 0, 0, 0, 0, 0, 18, 6050, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.043478260869565216, 0.25, 'gini = 0.017\nsamples = 1139\nvalue = [0, 0, 0, 0, 0, 0, 16, 1806, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.08695652173913043, 0.25, 'NMHC <= 0.265\ngini = 0.001\nsamples = 2656\nvalue = [0, 0, 0, 0, 0, 0, 2, 4244, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.06521739130434782, 0.08333333333333333, 'gini = 0.0\nsamples = 2415\nvalue = [0, 0, 0, 0, 0, 0, 0, 3877, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.10869565217391304, 0.08333333333333333, 'gini = 0.011\nsamples = 241\nvalue = [0, 0, 0, 0, 0, 0, 2, 367, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = h'),
Text(0.30434782608695654, 0.5833333333333334, 'CO <= 0.95\ngini = 0.91\nsamples = 42757\nvalue = [6228, 43, 6216, 6072, 6073, 28, 10, 17, 6146, 6\n28, 6077, 25, 30, 28, 6135, 57, 6091, 9, 6081, 21\n6305, 6041, 27]\nlass = v'),
Text(0.21739130434782608, 0.4166666666666667, 'O_3 <= 1.5\ngini = 0.801\nsamples = 18709\nvalue = [5952, 11, 0, 5959, 0, 20, 1, 0, 5990, 4, 0, 5876\n0, 0, 0, 0, 0, 0, 0, 5763, 20, 0, 0, 0]\nlass = i'),
Text(0.17391304347826086, 0.25, 'NO_2 <= 46.5\ngini = 0.064\nsamples = 3835\nvalue = [5952, 0, 0, 7, 0, 0, 0, 0, 123, 4, 0, 27, 0\n0, 0, 0, 0, 0, 0, 22, 20, 0, 0, 0]\nlass = a'),
Text(0.15217391304347827, 0.08333333333333333, 'gini = 0.02\nsamples = 2376\nvalue = [3755, 0, 0, 5, 0, 0, 0, 0, 8, 4, 0, 10, 0, 0\n0, 0, 0, 0, 0, 12, 0, 0, 0]\nlass = a'),
Text(0.1956521739130435, 0.08333333333333333, 'gini = 0.132\nsamples = 1459\nvalue = [2197, 0, 0, 2, 0, 0, 0, 0, 115, 0, 0, 17, 0\n0, 0, 0, 0, 0, 22, 8, 0, 0, 0]\nlass = a'),
Text(0.2608695652173913, 0.25, 'CO <= 0.15\ngini = 0.751\nsamples = 14874\nvalue = [0, 11, 0, 5952, 0, 20, 1, 0, 5867, 0, 0, 5849\n0, 0, 0, 0, 0, 0, 5741, 0, 0, 0, 0]\nlass = d'),
Text(0.2391304347826087, 0.08333333333333333, 'gini = 0.603\nsamples = 886\nvalue = [0, 3, 0, 36, 0, 0, 0, 0, 183, 0, 0, 435, 0\n0, 0, 0, 0, 750, 0, 0, 0, 0]\nlass = t'),
Text(0.2826086956521739, 0.08333333333333333, 'gini = 0.75\nsamples = 13988\nvalue = [0, 8, 0, 5916, 0, 20, 1, 0, 5684, 0, 0, 5414\n0, 0, 0, 0, 0, 0, 4991, 0, 0, 0, 0]\nlass = d'),
Text(0.391304347826087, 0.4166666666666667, 'NO <= 2.5\ngini = 0.845\nsamples = 24048\nvalue = [276, 32, 6216, 113, 6073, 8, 9, 17, 156, 2, 28\n201, 25, 30, 28, 6135, 57, 6091, 9, 318, 1, 6305\n6041, 27]\nlass = v'),
Text(0.34782608695652173, 0.25, 'BEN <= 0.95\ngini = 0.786\nsamples = 10298\nvalue = [13, 29, 1733, 8, 2225, 8, 9, 8, 3, 2, 13, 10\n13, 7, 3, 3859, 50, 163, 3, 32, 1, 4855, 3311\n24]\nlass = v'),
Text(0.32608695652173914, 0.08333333333333333, 'gini = 0.013\nsamples = 662\nvalue = [0, 2, 1050, 0, 0, 0, 0, 0, 0, 0, 5, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nlass = c'),
Text(0.3695652173913043, 0.08333333333333333, 'gini = 0.766\nsamples = 9636\nvalue = [13, 27, 683, 8, 2225, 8, 9, 8, 3, 2, 8, 10, 13\n7, 3, 3859, 50, 163, 3, 32, 1, 4855, 3311,
```



```
24]\nclasse = v'),
  Text(0.43478260869565216, 0.25, 'O_3 <= 1.5\ngini = 0.821\nsamples = 13750\nvalue = [263
, 3, 4483, 105, 3848, 0, 0, 9, 153, 0, 15\n191, 12, 23, 25, 2276, 7, 5928, 6, 286, 0, 145
0\n2730, 3]\nclasse = r'),
  Text(0.41304347826086957, 0.08333333333333333, 'gini = 0.29\nsamples = 3363\nvalue = [26
3, 3, 4483, 5, 18, 0, 0, 4, 44, 0, 15, 7, 12\n23, 25, 319, 7, 17, 6, 12, 0, 65, 15, 0]\ncl
asse = c'),
  Text(0.45652173913043476, 0.08333333333333333, 'gini = 0.768\nsamples = 10387\nvalue = [
0, 0, 0, 100, 3830, 0, 0, 5, 109, 0, 0, 184, 0\n0, 0, 1957, 0, 5911, 0, 274, 0, 1385, 271
5, 3]\nclasse = r'),
  Text(0.7934782608695652, 0.75, 'TCH <= 1.075\ngini = 0.917\nsamples = 46171\nvalue = [0,
5990, 0, 0, 0, 6082, 6151, 0, 0, 6120, 6126\n0, 5867, 6056, 6116, 0, 5907, 0, 6027, 0, 61
96, 0\n0, 6237]\nclasse = x'),
  Text(0.6521739130434783, 0.58333333333333334, 'SO_2 <= 1.5\ngini = 0.901\nsamples = 38603
\nvalue = [0, 5990, 0, 0, 0, 6082, 132, 0, 0, 6120, 6126, 0\n5867, 6056, 6116, 0, 5907, 0
, 62, 0, 6196, 0, 0\n6237]\nclasse = x'),
  Text(0.5652173913043478, 0.41666666666666667, 'NO_2 <= 28.5\ngini = 0.806\nsamples = 1818
7\nvalue = [0, 14, 0, 0, 0, 32, 24, 0, 0, 3296, 4, 0, 915\n6056, 6116, 0, 5907, 0, 62, 0,
21, 0, 0, 6237]\nclasse = x'),
  Text(0.5217391304347826, 0.25, 'NO <= 4.5\ngini = 0.809\nsamples = 9731\nvalue = [0, 3,
0, 0, 0, 13, 24, 0, 0, 2475, 3, 0, 639\n3425, 2779, 0, 2129, 0, 43, 0, 18, 0, 0, 3866]\ncl
asse = x'),
  Text(0.5, 0.08333333333333333, 'gini = 0.795\nsamples = 7232\nvalue = [0, 1, 0, 0, 0, 12
, 23, 0, 0, 1685, 1, 0, 559\n2857, 1555, 0, 1369, 0, 28, 0, 13, 0, 0, 3407]\nclasse = x'),
  Text(0.5434782608695652, 0.08333333333333333, 'gini = 0.788\nsamples = 2499\nvalue = [0,
2, 0, 0, 0, 1, 1, 0, 0, 790, 2, 0, 80\n568, 1224, 0, 760, 0, 15, 0, 5, 0, 0, 459]\nclasse =
o'),
  Text(0.6086956521739131, 0.25, 'O_3 <= 1.5\ngini = 0.78\nsamples = 8456\nvalue = [0, 11,
0, 0, 0, 19, 0, 0, 0, 821, 1, 0, 276\n2631, 3337, 0, 3778, 0, 19, 0, 3, 0, 0, 2371]\nclasse
= q'),
  Text(0.5869565217391305, 0.08333333333333333, 'gini = 0.723\nsamples = 6968\nvalue = [0,
0, 0, 0, 0, 3, 0, 0, 0, 821, 1, 0, 276\n2631, 3337, 0, 3778, 0, 19, 0, 3, 0, 0, 54]\nclasse
= q'),
  Text(0.6304347826086957, 0.08333333333333333, 'gini = 0.023\nsamples = 1488\nvalue = [0,
11, 0, 0, 0, 16, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 2317]\nclasse = x'),
  Text(0.7391304347826086, 0.41666666666666667, 'O_3 <= 1.5\ngini = 0.826\nsamples = 20416\
nvalue = [0, 5976, 0, 0, 0, 6050, 108, 0, 0, 2824, 6122, 0\n4952, 0, 0, 0, 0, 0, 0, 0, 0, 61
5, 0, 0, 0]\nclasse = u'),
  Text(0.6956521739130435, 0.25, 'EBE <= 0.95\ngini = 0.735\nsamples = 12767\nvalue = [0,
4, 0, 0, 0, 112, 2, 0, 0, 2824, 6122, 0\n4952, 0, 0, 0, 0, 0, 0, 0, 0, 6175, 0, 0, 0]\nclasse
= u'),
  Text(0.6739130434782609, 0.08333333333333333, 'gini = 0.041\nsamples = 2871\nvalue = [0,
3, 0, 0, 0, 88, 2, 0, 0, 0, 4366, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nclasse = k'),
  Text(0.717391304347826, 0.08333333333333333, 'gini = 0.702\nsamples = 9896\nvalue = [0,
1, 0, 0, 0, 24, 0, 0, 0, 2824, 1756, 0\n4952, 0, 0, 0, 0, 0, 0, 0, 0, 6175, 0, 0, 0]\nclasse
= u'),
  Text(0.782608695652174, 0.25, 'PM10 <= 15.5\ngini = 0.509\nsamples = 7649\nvalue = [0, 5
972, 0, 0, 0, 5938, 106, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nclasse = b'),
  Text(0.7608695652173914, 0.08333333333333333, 'gini = 0.503\nsamples = 3710\nvalue = [0,
2600, 0, 0, 0, 3209, 51, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nclasse = f'),
  Text(0.8043478260869565, 0.08333333333333333, 'gini = 0.503\nsamples = 3939\nvalue = [0,
3372, 0, 0, 0, 2729, 55, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nclasse = b'),
  Text(0.9347826086956522, 0.58333333333333334, 'O_3 <= 1.5\ngini = 0.5\nsamples = 7568\nva
lue = [0, 0, 0, 0, 0, 0, 6019, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 5965, 0, 0, 0, 0, 0]\nclas
s = g'),
  Text(0.9130434782608695, 0.41666666666666667, 'PM10 <= 15.5\ngini = 0.003\nsamples = 3788
\nvalue = [0, 0, 0, 0, 0, 0, 10, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 5965, 0, 0, 0, 0, 0]\ncl
asse = s'),
  Text(0.8695652173913043, 0.25, 'NO_2 <= 2.5\ngini = 0.001\nsamples = 2176\nvalue = [0, 0
, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 3399, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.8478260869565217, 0.08333333333333333, 'gini = 0.095\nsamples = 15\nvalue = [0, 0
, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 19, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.8913043478260869, 0.08333333333333333, 'gini = 0.001\nsamples = 2161\nvalue = [0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 3380, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.9565217391304348, 0.25, 'TCH <= 1.355\ngini = 0.006\nsamples = 1612\nvalue = [0,
0, 0, 0, 0, 0, 8, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 2566, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.9347826086956522, 0.08333333333333333, 'gini = 0.32\nsamples = 20\nvalue = [0, 0,
0, 0, 0, 0, 8, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 32, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.9782608695652174, 0.08333333333333333, 'gini = 0.0\nsamples = 1592\nvalue = [0, 0
, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 2534, 0, 0, 0, 0, 0]\nclasse = s'),
  Text(0.9565217391304348, 0.41666666666666667, 'gini = 0.0\nsamples = 3780\nvalue = [0, 0,
0, 0, 0, 0, 6009, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0]\nclasse = q')]
```



## Conclusion

## Accuracy

In [62]:

```
print("Linear Regression:", lr.score(x_test, y_test))
print("Ridge Regression:", rr.score(x_test, y_test))
print("Lasso Regression", la.score(x_test, y_test))
print("ElasticNet Regression:", en.score(x_test, y_test))
print("Logistic Regression:", logr.score(fs, target_vector))
print("Random Forest:", grid_search.best_score_)
```

Linear Regression: 0.30543236748624947  
Ridge Regression: 0.3054128646053008  
Lasso Regression 0.04531888865719813  
ElasticNet Regression: 0.1609692866422956  
Logistic Regression: 0.9  
Random Forest: 0.7496052165863487

## Logistic Regression is suitable for this dataset