

20104169 - SUMESH R

Importing Libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from google.colab import drive
drive.mount('/content/drive')
df=pd.read_csv("/content/drive/MyDrive/mydatasets/csvs_per_year/madrid_2017.csv")
df
```

Mounted at /content/drive

		date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx
0_3 \										
0	NaN	2017-06-01 01:00:00	NaN	NaN	0.3	NaN	NaN	4.0	38.0	NaN
1	71.0	2017-06-01 01:00:00	0.6	NaN	0.3	0.4	0.08	3.0	39.0	NaN
2	NaN	2017-06-01 01:00:00	0.2	NaN	NaN	0.1	NaN	1.0	14.0	NaN
3	91.0	2017-06-01 01:00:00	NaN	NaN	0.2	NaN	NaN	1.0	9.0	NaN
4	69.0	2017-06-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	19.0	NaN
...	
...										
210115	65.0	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	27.0	NaN
210116	NaN	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	14.0	NaN
210117	83.0	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	4.0	NaN
210118	78.0	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	11.0	NaN
210119	77.0	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	14.0	NaN
		PM10	PM25	SO_2	TCH	TOL	station			
0		NaN	NaN	5.0	NaN	NaN	28079004			
1		22.0	9.0	7.0	1.4	2.9	28079008			
2		NaN	NaN	NaN	NaN	0.9	28079011			

```

3      NaN    NaN    NaN    NaN    NaN    28079016
4      NaN    NaN    2.0    NaN    NaN    28079017
...
210115    NaN    NaN    NaN    NaN    NaN    28079056
210116    73.0    NaN    7.0    NaN    NaN    28079057
210117    NaN    NaN    NaN    NaN    NaN    28079058
210118    NaN    NaN    NaN    NaN    NaN    28079059
210119    60.0    NaN    NaN    NaN    NaN    28079060

```

```
[210120 rows x 16 columns]
```

Data Cleaning and Data Preprocessing

```
df=df.dropna()
```

```
df.columns
```

```
Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx',
      'O_3',
      'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 4127 entries, 87457 to 158286
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        4127 non-null   object
 1   BEN         4127 non-null   float64
 2   CH4         4127 non-null   float64
 3   CO          4127 non-null   float64
 4   EBE         4127 non-null   float64
 5   NMHC        4127 non-null   float64
 6   NO          4127 non-null   float64
 7   NO_2        4127 non-null   float64
 8   NOx         4127 non-null   float64
 9   O_3         4127 non-null   float64
10  PM10        4127 non-null   float64
11  PM25        4127 non-null   float64
12  SO_2        4127 non-null   float64
13  TCH         4127 non-null   float64
14  TOL         4127 non-null   float64
15  station     4127 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 548.1+ KB

```

```
data=df[['CO' , 'station']]
data
```


	CO	station
87457	0.3	28079008
87462	0.2	28079024
87481	0.2	28079008
87486	0.2	28079024
87505	0.2	28079008
...
158238	0.2	28079024
158257	0.3	28079008
158262	0.2	28079024
158281	0.2	28079008
158286	0.2	28079024

```
[4127 rows x 2 columns]
```

Line chart

```
data.plot.line(subplots=True)
```


```
array([<Axes: >, <Axes: >], dtype=object)
```



Line chart

```
data.plot.line()
```

<Axes: >




Bar chart

```
b=data[0:50]
```

```
b.plot.bar()
```


<Axes: >



Histogram

```
data.plot.hist()
```


```
<Axes: ylabel='Frequency'>
```



Area chart

```
data.plot.area()
```


<Axes: >



Box chart

```
data.plot.box()
```


```
<Axes: >
```



Pie chart

```
b.plot.pie(y='station' )
```


```
<Axes: ylabel='station'>
```

Scatter chart

```
data.plot.scatter(x='CO' ,y='station')
```

```
<Axes: xlabel='CO', ylabel='station'>
```



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 4127 entries, 87457 to 158286
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	date	4127 non-null	object
1	BEN	4127 non-null	float64
2	CH4	4127 non-null	float64
3	CO	4127 non-null	float64
4	EBE	4127 non-null	float64
5	NMHC	4127 non-null	float64
6	NO	4127 non-null	float64
7	NO_2	4127 non-null	float64
8	NOx	4127 non-null	float64
9	O_3	4127 non-null	float64

```

10 PM10      4127 non-null float64
11 PM25      4127 non-null float64
12 SO_2      4127 non-null float64
13 TCH       4127 non-null float64
14 TOL       4127 non-null float64
15 station   4127 non-null int64
dtypes: float64(14), int64(1), object(1)
memory usage: 548.1+ KB

```

```
df.describe()
```

	BEN	CH4	CO	EBE	NMHC
\					
count	4127.000000	4127.000000	4127.000000	4127.000000	4127.000000
mean	0.919918	1.323732	0.417858	0.578168	0.097269
std	1.123078	0.215742	0.342871	0.962000	0.094035
min	0.100000	1.100000	0.100000	0.100000	0.000000
25%	0.300000	1.180000	0.200000	0.100000	0.050000
50%	0.600000	1.270000	0.300000	0.300000	0.080000
75%	1.100000	1.400000	0.500000	0.700000	0.110000
max	19.600000	3.630000	4.900000	16.700001	1.420000

	NO	NO_2	NOx	O_3	PM10
\					
count	4127.000000	4127.000000	4127.000000	4127.000000	4127.000000
mean	41.785316	58.069057	122.125515	28.716501	17.582021
std	71.118499	38.974112	142.828344	25.304909	12.735860
min	1.000000	1.000000	2.000000	1.000000	1.000000
25%	3.000000	30.000000	37.000000	6.000000	8.000000
50%	16.000000	54.000000	80.000000	22.000000	14.000000
75%	50.000000	78.000000	153.000000	46.000000	25.000000
max	879.000000	349.000000	1681.000000	140.000000	80.000000

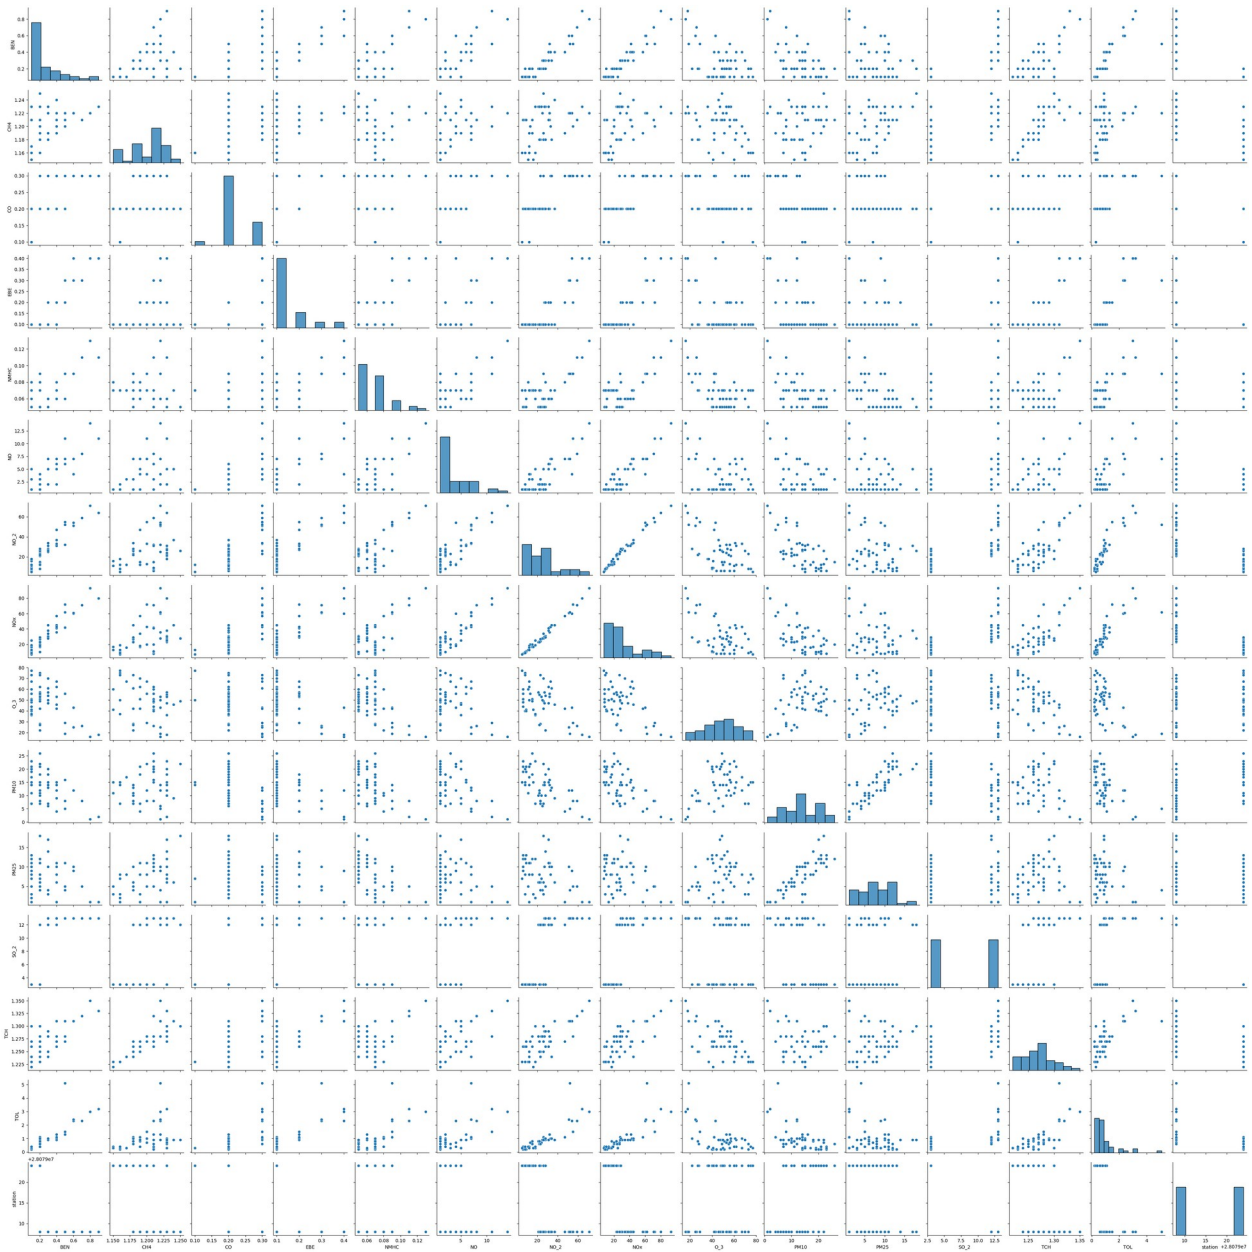
	PM25	SO_2	TCH	TOL
station				

count	4127.000000	4127.000000	4127.000000	4127.000000
4.127000e+03				
mean	10.942816	5.689120	1.420417	4.162830
2.807902e+07				
std	8.511526	3.848442	0.261857	5.689394
8.000152e+00				
min	1.000000	1.000000	1.100000	0.100000
2.807901e+07				
25%	5.000000	3.000000	1.260000	1.000000
2.807901e+07				
50%	9.000000	4.000000	1.370000	2.500000
2.807901e+07				
75%	15.000000	7.000000	1.480000	5.200000
2.807902e+07				
max	58.000000	32.000000	3.700000	84.800003
2.807902e+07				

EDA AND VISUALIZATION

```
sns.pairplot(df[0:50])
```

```
<seaborn.axisgrid.PairGrid at 0x7981b67df430>
```



```
sns.distplot(df['station'])
```

```
<ipython-input-19-6e2460d4583e>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['station'])
```

<Axes: xlabel='station', ylabel='Density'>




```
sns.heatmap(df.corr())
```

<ipython-input-20-aa4f4450a243>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr())
```

<Axes: >



TO TRAIN THE MODEL AND MODEL BUILDING

```
x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

LinearRegression()

```
lr.intercept_
```


28079042.226563875

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

	Co-efficient
BEN	0.043547
CO	-5.098540
EBE	-1.826949
NMHC	22.952181
NO	0.053629
NO_2	-0.183132
O_3	-0.092828
PM10	0.456248
PM25	-0.203910
SO_2	-0.260461
TCH	-13.705300
TOL	0.200241

```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
<matplotlib.collections.PathCollection at 0x79819dee8cd0>
```



ACCURACY

```
lr.score(x_test,y_test)
0.6415831920868991
lr.score(x_train,y_train)
0.632157725410589
```

Ridge and Lasso

```
from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
Ridge(alpha=10)
```

Accuracy(Ridge)

```
rr.score(x_test,y_test)
0.6253250346804637
rr.score(x_train,y_train)
0.6249671707210255
la=Lasso(alpha=10)
la.fit(x_train,y_train)
Lasso(alpha=10)
la.score(x_train,y_train)
0.40470624691546153
```

Accuracy(Lasso)

```
la.score(x_test,y_test)
0.40569400390890953
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```

ElasticNet()
en.coef_
array([-0.          , -0.          , -0.          ,  0.          ,
        0.03093261,
        -0.20550351, -0.08746041,  0.56446666, -0.41952604, -
        0.28904518,
        -0.          ,  0.          ])
en.intercept_
28079025.37399155
prediction=en.predict(x_test)
en.score(x_test,y_test)
0.5044230795721029

```

Evaluation Metrics

```

from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))

4.814433281391835
31.697067819145467
5.63001490398964

```

Logistic Regression

```

from sklearn.linear_model import LogisticRegression

feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3',
'PM10', 'PM25',
'SO_2', 'TCH', 'TOL']]
target_vector=df['station']

feature_matrix.shape
(4127, 12)

target_vector.shape
(4127,)

from sklearn.preprocessing import StandardScaler

```

```

fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)

LogisticRegression(max_iter=10000)
observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
prediction=logr.predict(observation)
print(prediction)

[28079008]

logr.classes_
array([28079008, 28079024])

logr.score(fs,target_vector)
0.9520232614489944

logr.predict_proba(observation)[0][0]
0.9999999999999389

logr.predict_proba(observation)
array([[1.00000000e+00, 6.10312359e-14]])

```

Random Forest

```

from sklearn.ensemble import RandomForestClassifier

rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)

RandomForestClassifier()

parameters={'max_depth':[1,2,3,4,5],
            'min_samples_leaf':[5,10,15,20,25],
            'n_estimators':[10,20,30,40,50]
}

from sklearn.model_selection import GridSearchCV
grid_search
=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)

```

```

GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                          'min_samples_leaf': [5, 10, 15, 20, 25],
                          'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')

grid_search.best_score_

0.9719529085872576

rfc_best=grid_search.best_estimator_

from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=
['a', 'b', 'c', 'd'], filled=True)

[Text(0.4107142857142857, 0.9166666666666666, 'TCH <= 1.345\ngini =
0.5\nsamples = 1806\nvalue = [1478, 1410]\nclass = a'),
Text(0.19285714285714287, 0.75, 'TCH <= 1.285\ngini = 0.225\nsamples
= 796\nvalue = [163, 1096]\nclass = b'),
Text(0.08571428571428572, 0.5833333333333334, 'SO_2 <= 10.5\ngini =
0.052\nsamples = 542\nvalue = [23, 834]\nclass = b'),
Text(0.05714285714285714, 0.4166666666666667, 'BEN <= 0.45\ngini =
0.017\nsamples = 531\nvalue = [7, 834]\nclass = b'),
Text(0.02857142857142857, 0.25, 'gini = 0.0\nsamples = 362\nvalue =
[0, 574]\nclass = b'),
Text(0.08571428571428572, 0.25, 'O_3 <= 36.0\ngini = 0.051\nsamples =
169\nvalue = [7, 260]\nclass = b'),
Text(0.05714285714285714, 0.08333333333333333, 'gini = 0.0\nsamples =
144\nvalue = [0, 231]\nclass = b'),
Text(0.11428571428571428, 0.08333333333333333, 'gini = 0.313\nsamples
= 25\nvalue = [7, 29]\nclass = b'),
Text(0.11428571428571428, 0.4166666666666667, 'gini = 0.0\nsamples =
11\nvalue = [16, 0]\nclass = a'),
Text(0.3, 0.5833333333333334, 'PM10 <= 7.5\ngini = 0.454\nsamples =
254\nvalue = [140, 262]\nclass = b'),
Text(0.22857142857142856, 0.4166666666666667, 'EBE <= 0.15\ngini =
0.089\nsamples = 40\nvalue = [61, 3]\nclass = a'),
Text(0.2, 0.25, 'NMHC <= 0.085\ngini = 0.245\nsamples = 13\nvalue =
[18, 3]\nclass = a'),
Text(0.17142857142857143, 0.08333333333333333, 'gini = 0.0\nsamples =
7\nvalue = [12, 0]\nclass = a'),
Text(0.22857142857142856, 0.08333333333333333, 'gini = 0.444\nsamples
= 6\nvalue = [6, 3]\nclass = a'),
Text(0.2571428571428571, 0.25, 'gini = 0.0\nsamples = 27\nvalue =
[43, 0]\nclass = a'),
Text(0.37142857142857144, 0.4166666666666667, 'O_3 <= 35.5\ngini =
0.358\nsamples = 214\nvalue = [79, 259]\nclass = b'),

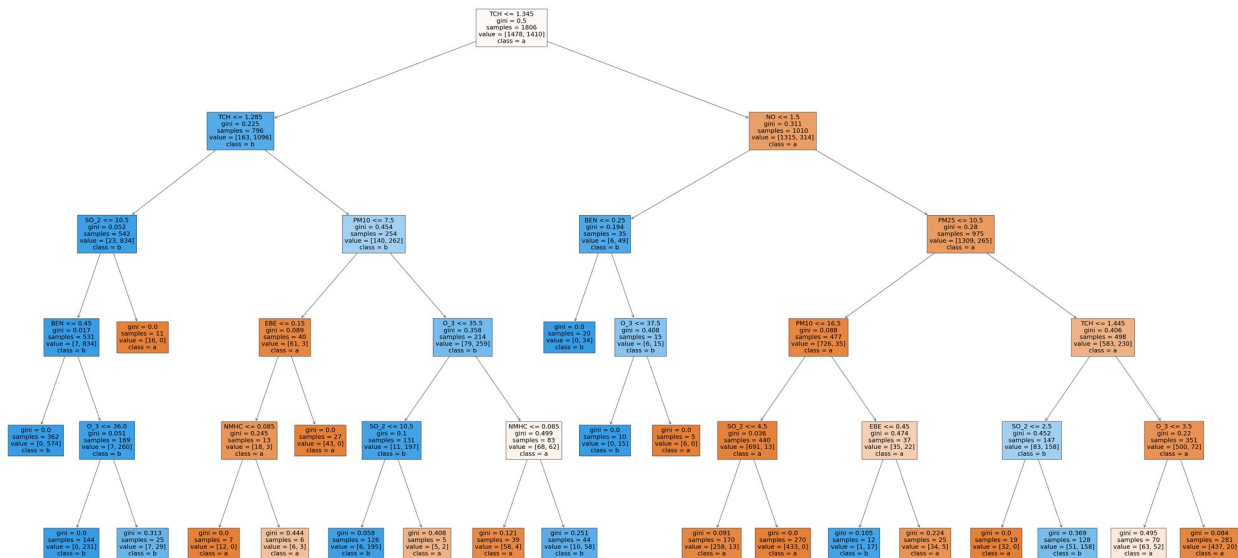
```

```
Text(0.3142857142857143, 0.25, 'SO_2 <= 10.5\ngini = 0.1\nsamples = 131\nvalue = [11, 197]\nclass = b'),
Text(0.2857142857142857, 0.08333333333333333, 'gini = 0.058\nsamples = 126\nvalue = [6, 195]\nclass = b'),
Text(0.34285714285714286, 0.08333333333333333, 'gini = 0.408\nsamples = 5\nvalue = [5, 2]\nclass = a'),
Text(0.42857142857142855, 0.25, 'NMHC <= 0.085\ngini = 0.499\nsamples = 83\nvalue = [68, 62]\nclass = a'),
Text(0.4, 0.08333333333333333, 'gini = 0.121\nsamples = 39\nvalue = [58, 4]\nclass = a'),
Text(0.45714285714285713, 0.08333333333333333, 'gini = 0.251\nsamples = 44\nvalue = [10, 58]\nclass = b'),
Text(0.6285714285714286, 0.75, 'NO <= 1.5\ngini = 0.311\nsamples = 1010\nvalue = [1315, 314]\nclass = a'),
Text(0.4857142857142857, 0.5833333333333334, 'BEN <= 0.25\ngini = 0.194\nsamples = 35\nvalue = [6, 49]\nclass = b'),
Text(0.45714285714285713, 0.4166666666666667, 'gini = 0.0\nsamples = 20\nvalue = [0, 34]\nclass = b'),
Text(0.5142857142857142, 0.4166666666666667, 'O_3 <= 37.5\ngini = 0.408\nsamples = 15\nvalue = [6, 15]\nclass = b'),
Text(0.4857142857142857, 0.25, 'gini = 0.0\nsamples = 10\nvalue = [0, 15]\nclass = b'),
Text(0.5428571428571428, 0.25, 'gini = 0.0\nsamples = 5\nvalue = [6, 0]\nclass = a'),
Text(0.7714285714285715, 0.5833333333333334, 'PM25 <= 10.5\ngini = 0.28\nsamples = 975\nvalue = [1309, 265]\nclass = a'),
Text(0.6571428571428571, 0.4166666666666667, 'PM10 <= 16.5\ngini = 0.088\nsamples = 477\nvalue = [726, 35]\nclass = a'),
Text(0.6, 0.25, 'SO_2 <= 4.5\ngini = 0.036\nsamples = 440\nvalue = [691, 13]\nclass = a'),
Text(0.5714285714285714, 0.08333333333333333, 'gini = 0.091\nsamples = 170\nvalue = [258, 13]\nclass = a'),
Text(0.6285714285714286, 0.08333333333333333, 'gini = 0.0\nsamples = 270\nvalue = [433, 0]\nclass = a'),
Text(0.7142857142857143, 0.25, 'EBE <= 0.45\ngini = 0.474\nsamples = 37\nvalue = [35, 22]\nclass = a'),
Text(0.6857142857142857, 0.08333333333333333, 'gini = 0.105\nsamples = 12\nvalue = [1, 17]\nclass = b'),
Text(0.7428571428571429, 0.08333333333333333, 'gini = 0.224\nsamples = 25\nvalue = [34, 5]\nclass = a'),
Text(0.8857142857142857, 0.4166666666666667, 'TCH <= 1.445\ngini = 0.406\nsamples = 498\nvalue = [583, 230]\nclass = a'),
Text(0.8285714285714286, 0.25, 'SO_2 <= 2.5\ngini = 0.452\nsamples = 147\nvalue = [83, 158]\nclass = b'),
Text(0.8, 0.08333333333333333, 'gini = 0.0\nsamples = 19\nvalue = [32, 0]\nclass = a'),
Text(0.8571428571428571, 0.08333333333333333, 'gini = 0.369\nsamples = 128\nvalue = [51, 158]\nclass = b'),
Text(0.9428571428571428, 0.25, 'O_3 <= 3.5\ngini = 0.22\nsamples =
```

```

351\nvalue = [500, 72]\nnclass = a'),
Text(0.9142857142857143, 0.08333333333333333, 'gini = 0.495\nsamples
= 70\nnvalue = [63, 52]\nnclass = a'),
Text(0.9714285714285714, 0.08333333333333333, 'gini = 0.084\nsamples
= 281\nnvalue = [437, 20]\nnclass = a')]

```



Conclusion

Accuracy

```

print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)

```

Linear Regression: 0.6415831920868991
 Ridge Regression: 0.6253250346804637
 Lasso Regression 0.40569400390890953
 ElasticNet Regression: 0.5044230795721029
 Logistic Regression: 0.9520232614489944
 Random Forest: 0.9719529085872576

Random Forest is suitable for this dataset